

▼ eCommerce Transactions Dataset EDA

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

#Load the all file:

customers=pd.read_csv('/content/Customers.csv')
products=pd.read_csv('/content/Products.csv')
transactions=pd.read_csv('/content/Transactions.csv')




#Check Data:
customers.head()
```



	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

Next steps: [Generate code with customers](#) [View recommended plots](#) [New interactive sheet](#)




```
products.head()
```



	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

Next steps: [Generate code with products](#) [View recommended plots](#) [New interactive sheet](#)

```
transactions.head()
```



	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	Price
0	T00001	C0199	P067	2024-08-25 12:38:23	1	300.68	300.68
1	T00112	C0146	P067	2024-05-27 22:23:54	1	300.68	300.68
2	T00166	C0127	P067	2024-04-25 07:38:55	1	300.68	300.68
3	T00272	C0087	P067	2024-03-26 22:55:37	2	601.36	300.68
4	T00363	C0070	P067	2024-03-21 15:10:10	3	902.04	300.68

Next steps: [Generate code with transactions](#) [View recommended plots](#) [New interactive sheet](#)

▼ #Data Understanding

```
#Check the data shape
print(customers.info())
print()
print(products.info())
print()
print(transactions.info())
print()
```

```

↳ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   CustomerID      200 non-null   object
 1   CustomerName    200 non-null   object
 2   Region          200 non-null   object
 3   SignupDate      200 non-null   object
dtypes: object(4)
memory usage: 6.4+ KB
None

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 4 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   ProductID       100 non-null   object
 1   ProductName     100 non-null   object
 2   Category        100 non-null   object
 3   Price           100 non-null   float64
dtypes: float64(1), object(3)
memory usage: 3.3+ KB
None

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 7 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   TransactionID    1000 non-null   object
 1   CustomerID       1000 non-null   object
 2   ProductID        1000 non-null   object
 3   TransactionDate   1000 non-null   object
 4   Quantity         1000 non-null   int64
 5   TotalValue       1000 non-null   float64
 6   Price            1000 non-null   float64
dtypes: float64(2), int64(1), object(4)
memory usage: 54.8+ KB
None

```

```

#Check missing Value
print(customers.isnull().sum())
print()
print(products.isnull().sum())
print()
print(transactions.isnull().sum())
print()

```

```

↳ CustomerID      0
CustomerName      0
Region            0
SignupDate        0
dtype: int64

ProductID         0
ProductName        0
Category          0
Price             0
dtype: int64

TransactionID      0
CustomerID         0
ProductID          0
TransactionDate    0
Quantity           0
TotalValue         0
Price             0
dtype: int64

```

```

#Check duplicate Value

```

```

print(customers.duplicated().sum())
print()
print(products.duplicated().sum())
print()
print(transactions.duplicated().sum())
print()

```

```
0
0
0
```

#Describe the transaction Table

```
round(transactions.describe()).T
```

	count	mean	std	min	25%	50%	75%	max
<b>Quantity</b>	1000.0	3.0	1.0	1.0	2.0	3.0	4.0	4.0
<b>TotalValue</b>	1000.0	690.0	493.0	16.0	295.0	589.0	1012.0	1991.0
<b>Price</b>	1000.0	273.0	141.0	16.0	148.0	300.0	404.0	498.0

## ✓ #Data Cleaning:

```
print(customers.head())
print()
print(products.head())
print()
```

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

#Remove the time from TransactionDate column of transaction Table:

```
transactions['TransactionDate'] = pd.to_datetime(transactions['TransactionDate'])
transactions['TransactionDate'] = transactions['TransactionDate'].dt.date
```

```
customers['SignupDate'] = pd.to_datetime(customers['SignupDate'])
transactions['TransactionDate'] = pd.to_datetime(transactions['TransactionDate'])
```

```
print(customers.head())
print()
print(transactions.head())
print()
```

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	\
0	T00001	C0199	P067	2024-08-25	1	300.68	
1	T00112	C0146	P067	2024-05-27	1	300.68	
2	T00166	C0127	P067	2024-04-25	1	300.68	
3	T00272	C0087	P067	2024-03-26	2	601.36	
4	T00363	C0070	P067	2024-03-21	3	902.04	

	Price
0	300.68
1	300.68
2	300.68
3	300.68
4	300.68

```
#Check the data shape
print(customers.info())
print()
print(products.info())
print()
print(transactions.info())
print()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   CustomerID      200 non-null   object
 1   CustomerName    200 non-null   object
 2   Region          200 non-null   object
 3   SignupDate      200 non-null   datetime64[ns]
dtypes: datetime64[ns](1), object(3)
memory usage: 6.4+ KB
None
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 4 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   ProductID       100 non-null   object
 1   ProductName     100 non-null   object
 2   Category        100 non-null   object
 3   Price           100 non-null   float64
dtypes: float64(1), object(3)
memory usage: 3.3+ KB
None
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 7 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   TransactionID    1000 non-null   object
 1   CustomerID       1000 non-null   object
 2   ProductID        1000 non-null   object
 3   TransactionDate   1000 non-null   datetime64[ns]
 4   Quantity         1000 non-null   int64
 5   TotalValue       1000 non-null   float64
 6   Price            1000 non-null   float64
dtypes: datetime64[ns](1), float64(2), int64(1), object(3)
memory usage: 54.8+ KB
None
```

## Exploratory Analysis (Lets analyze all the file indivisual)

### customers file:

```
customers.head()
```

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

Next steps:

[Generate code with customers](#)[View recommended plots](#)[New interactive sheet](#)

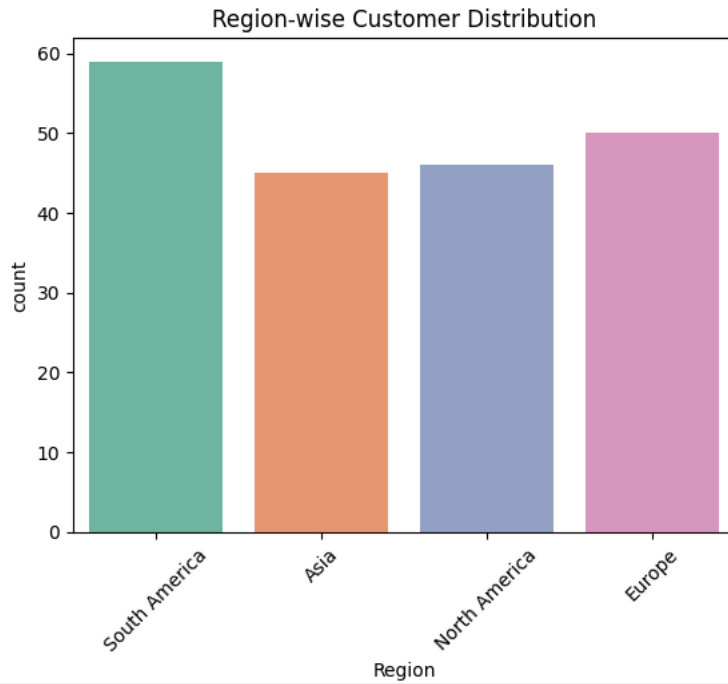
```
#Distributed region wise:
```

```
sns.countplot(data=customers, x='Region', palette='Set2')
plt.xticks(rotation=45)
plt.title("Region-wise Customer Distribution")
plt.show()
```

<ipython-input-65-e2c8c2e5aadd>:3: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set

```
sns.countplot(data=customers, x='Region', palette='Set2')
```



```
#Analysis by Signup date
customers['SignupMonth'] = customers['SignupDate'].dt.to_period('M')
monthly_signups = customers.groupby('SignupMonth').size()
```

```
monthly_signups.plot(kind='bar', figsize=(12, 6), color='skyblue')
plt.title("Monthly Signups")
plt.xlabel("Month")
plt.ylabel("Number of Customers")
plt.show()
```



▼ From Product File:


```
#Category_wise product count:
```

```
sns.countplot(data=products, x='Category', palette='Set2')  
plt.title("Product Count by Category")  
plt.xticks(rotation=45)  
plt.show()
```

 <ipython-input-54-2fad0c645677>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set

```
sns.countplot(data=products, x='Category', palette='Set2')
```



```
#price Distribution:
```

```
sns.histplot(products['Price'], bins=20, kde=True, color='orange')  
plt.title("Product Price Distribution")  
plt.show()
```



## Transactions File

```
transactions.head()
```

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	Price	
	0	T00001	C0199	P067	2024-08-25	1	300.68	300.68
	1	T00112	C0146	P067	2024-05-27	1	300.68	300.68
	2	T00166	C0127	P067	2024-04-25	1	300.68	300.68
	3	T00272	C0087	P067	2024-03-26	2	601.36	300.68
	4	T00363	C0070	P067	2024-03-21	3	902.04	300.68

Next steps:

[Generate code with transactions](#)[View recommended plots](#)[New interactive sheet](#)

#calculate Total revenue:

```
total_revenue = transactions['TotalValue'].sum()
print(f"Total Revenue: {total_revenue}")
```

Total Revenue: 689995.56

# Top 10 customers by revenue:

```
top_customers = transactions.groupby('CustomerID')['TotalValue'].sum().sort_values(ascending=False).head(10)
print(top_customers)
```

CustomerID

C0141	10673.87
C0054	8040.39
C0065	7663.70
C0156	7634.45
C0082	7572.91
C0188	7111.32
C0059	7073.28
C0028	6819.57
C0099	6715.72
C0165	6708.10

Name: TotalValue, dtype: float64

# Top 10 products by sales:

```
top_products = transactions.groupby('ProductID')['Quantity'].sum().sort_values(ascending=False).head(10)
print(top_products)
```

ProductID


P059	46
P054	46
P029	45
P079	43
P061	43
P057	43
P048	43
P062	39
P020	38
P028	38

Name: Quantity, dtype: int64

#Monthly revenue trends:

```
transactions['TransactionMonth'] = transactions['TransactionDate'].dt.to_period('M')
monthly_revenue = transactions.groupby('TransactionMonth')['TotalValue'].sum()

monthly_revenue.plot(kind='line', figsize=(12, 6), marker='o', color='green')
plt.title("Monthly Revenue Trends")
plt.xlabel("Month")
plt.ylabel("Revenue")
plt.show()
```



## Business Insights:

### 1. Customer Demographics and Signup Trends:

A. South America has the highest number of customers compared to other regions. This indicates that South America is a key market for the business.

B. Most customers signed up during February, March, and April, whereas January consistently shows the lowest number of new signups.

### 2. Product Category and Pricing Trends:

A. Products across categories have almost equal distribution, with a slight dip in the Home Decor category. This might indicate either lower demand or limited product variety in this segment.

B. Price distribution highlights two key ranges: 100–200 and 400–500, showing these are the most popular pricing brackets among customers.

### 3. Revenue and Customer Trends:

A. The total revenue generated is 689,995.56, with *CustomerID : C0141* being the top contributor, generating 10,673 in revenue.

B. The top-selling product is P059, sold 46 times, indicating strong demand for this product.

### 4. Monthly Revenue Trends:

A. In December 2023, revenue was critically low (below 1,000), but it showed a sharp recovery to over 6,000 in January 2024.

B. Post-recovery, the monthly revenue stabilized in the range of 5,000–7,000 for the rest of the year, showing consistent performance after an initial spike.

## ✓ TASK 2 : Looklike Model

```
#import sklearn library
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.preprocessing import MinMaxScaler
```

```
# visit the Data:
```

```
print(customers.head())
print()
```



```
print(products.head())
print()
print(transactions.head())
```

	CustomerID	CustomerName	Region	SignupDate	SignupMonth
0	C0001	Lawrence Carroll	South America	2022-07-10	2022-07
1	C0002	Elizabeth Lutz	Asia	2022-02-13	2022-02
2	C0003	Michael Rivera	South America	2024-03-07	2024-03
3	C0004	Kathleen Rodriguez	South America	2022-10-09	2022-10
4	C0005	Laura Weber	Asia	2022-08-15	2022-08

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	\
0	T00001	C0199	P067	2024-08-25	1	300.68	
1	T00112	C0146	P067	2024-05-27	1	300.68	
2	T00166	C0127	P067	2024-04-25	1	300.68	
3	T00272	C0087	P067	2024-03-26	2	601.36	
4	T00363	C0070	P067	2024-03-21	3	902.04	

	Price	TransactionMonth
0	300.68	2024-08
1	300.68	2024-05
2	300.68	2024-04
3	300.68	2024-03
4	300.68	2024-03

```
# Merge 3 Dataset
# Merge transactions with products
transactions = transactions.merge(products, on='ProductID', how='left')

# Merge transactions with customers
data = transactions.merge(customers, on='CustomerID', how='left')

# Check the final dataset
data.head()
```

TransactionID	CustomerID	ProductID	TransactionDate	Quantity	TotalValue	Price_x	TransactionMonth	ProductName_x	Category
T00001	C0199	P067	2024-08-25	1	300.68	300.68	2024-08	ComfortLiving Bluetooth Speaker	Electronics
T00112	C0146	P067	2024-05-27	1	300.68	300.68	2024-05	ComfortLiving Bluetooth Speaker	Electronics
T00166	C0127	P067	2024-04-25	1	300.68	300.68	2024-04	ComfortLiving Bluetooth Speaker	Electronics
T00272	C0087	P067	2024-03-26	2	601.36	300.68	2024-03	ComfortLiving Bluetooth Speaker	Electronics
T00363	C0070	P067	2024-03-21	3	902.04	300.68	2024-03	ComfortLiving Bluetooth Speaker	Electronics

Next steps:

Generate code with data

View recommended plots

New interactive sheet

```
# Aggregate transaction data
customer_agg = data.groupby('CustomerID').agg({'TotalValue': 'sum', 'Quantity': 'sum', 'Price': 'mean'}).reset_index()

# Add region encoding
region_encoding = pd.get_dummies(customers[['CustomerID', 'Region']], columns=['Region'])
customer_agg = customer_agg.merge(region_encoding, on='CustomerID', how='left')

# Check aggregated features
customer_agg.head()
```

	CustomerID	TotalValue	Quantity	Price	Region_Asia	Region_Europe	Region_North America	Region_South America
0	C0001	3354.52	12	278.334000	False	False	False	True
1	C0002	1862.74	10	208.920000	True	False	False	False
2	C0003	2725.38	14	195.707500	False	False	False	True
3	C0004	5354.88	23	240.636250	False	False	False	True
4	C0005	2034.24	7	291.603333	True	False	False	False

Next steps:

[Generate code with customer\\_agg](#)[View recommended plots](#)[New interactive sheet](#)

## Normalize the Features

```
# Normalize numerical features
scaler = MinMaxScaler()
features = customer_agg.drop('CustomerID', axis=1) # Exclude CustomerID for scaling
normalized_features = scaler.fit_transform(features)

# Convert back to DataFrame
normalized_features = pd.DataFrame(normalized_features, columns=features.columns, index=customer_agg['CustomerID'])
normalized_features.head()
```

	TotalValue	Quantity	Price	Region_Asia	Region_Europe	Region_North America	Region_South America
CustomerID							
C0001	0.308942	0.354839	0.519414	0.0	0.0	0.0	1.0
C0002	0.168095	0.290323	0.367384	1.0	0.0	0.0	0.0
C0003	0.249541	0.419355	0.338446	0.0	0.0	0.0	1.0
C0004	0.497806	0.709677	0.436848	0.0	0.0	0.0	1.0
C0005	0.184287	0.193548	0.548476	1.0	0.0	0.0	0.0

Next steps:

[Generate code with normalized\\_features](#)[View recommended plots](#)[New interactive sheet](#)

## Calculate similirity score (Build the lookalike model)

```
# Compute pairwise cosine similarity
similarity_matrix = cosine_similarity(normalized_features)

# Convert to DataFrame for easy manipulation
similarity_df = pd.DataFrame(similarity_matrix, index=customer_agg['CustomerID'], columns=customer_agg['CustomerID'])
similarity_df.head()
```

CustomerID	C0001	C0002	C0003	C0004	C0005	C0006	C0007	C0008	C0009	C0010	...	C0191	C0192
CustomerID													
C0001	1.000000	0.253518	0.986950	0.959271	0.286969	0.995359	0.330221	0.348553	0.241389	0.224175	...	0.999278	0.989804
C0002	0.253518	1.000000	0.221730	0.289222	0.984924	0.277859	0.969086	0.272393	0.185625	0.178380	...	0.247901	0.224718
C0003	0.986950	0.221730	1.000000	0.968405	0.229577	0.968467	0.262941	0.326945	0.176962	0.211370	...	0.984116	0.970527
C0004	0.959271	0.289222	0.968405	1.000000	0.287091	0.953520	0.328818	0.447073	0.205789	0.284352	...	0.948407	0.910928
C0005	0.286969	0.984924	0.229577	0.287091	1.000000	0.322382	0.995295	0.274389	0.248972	0.179873	...	0.285261	0.271305

5 rows × 199 columns

## Generate recoomendations

```
# Function to get top 3 lookalikes for each customer
def get_top_lookalikes(similarity_df, top_n=3):
    lookalike_list = []
    for customer_id in similarity_df.index:
        # Sort customers by similarity score, skip self (iloc[1:top_n+1])
        similar_customers = similarity_df.loc[customer_id].sort_values(ascending=False).iloc[1:top_n+1]
        # Add results to the list
```

```

    lookalike_list.append(
        [customer_id] + list(similar_customers.index) + list(similar_customers.values)
    )
    return lookalike_list

# Get top 3 lookalikes for all customers
lookalike_list = get_top_lookalikes(similarity_df)

# Convert the list into a DataFrame
lookalike_df = pd.DataFrame(
    lookalike_list,
    columns=['CustomerID', 'Lookalike1', 'Lookalike2', 'Lookalike3', 'Score1', 'Score2', 'Score3']
)

# Save the DataFrame as a CSV file
lookalike_df.to_csv('Lookalike.csv', index=False)

# Show the first few rows
lookalike_df.head()

```

	CustomerID	Lookalike1	Lookalike2	Lookalike3	Score1	Score2	Score3
0	C0001	C0137	C0191	C0011	0.999479	0.999278	0.999276
1	C0002	C0088	C0142	C0027	0.998965	0.998596	0.996510
2	C0003	C0190	C0147	C0174	0.998906	0.997914	0.996201
3	C0004	C0113	C0169	C0012	0.999112	0.995818	0.994859
4	C0005	C0186	C0140	C0146	0.998743	0.998450	0.996784

Next steps:

[Generate code with lookalike\\_df](#)[View recommended plots](#)[New interactive sheet](#)

```

#Top 3 similar customers and their scores for all customers
lookalike_df.head(3)

```

	CustomerID	Lookalike1	Lookalike2	Lookalike3	Score1	Score2	Score3
0	C0001	C0137	C0191	C0011	0.999479	0.999278	0.999276
1	C0002	C0088	C0142	C0027	0.998965	0.998596	0.996510
2	C0003	C0190	C0147	C0174	0.998906	0.997914	0.996201

Next steps:

[Generate code with lookalike\\_df](#)[View recommended plots](#)[New interactive sheet](#)

## Task 3: Customer Segmentation / Clustering

To group customers into segments based on their profiles and transaction history, enabling better targeting and personalized marketing.

```

# Merge Customers.csv and Transactions.csv
merged_df = pd.merge(transactions, customers, on='CustomerID', how='left')

# Aggregate transaction-level data for each customer
customer_agg = merged_df.groupby('CustomerID').agg({
    'TotalValue': 'sum',          # Total revenue per customer
    'TransactionID': 'count',     # Total number of transactions
    'Quantity': 'mean',          # Average basket size
}).rename(columns={
    'TotalValue': 'TotalRevenue',
    'TransactionID': 'TotalTransactions',
    'Quantity': 'AvgBasketSize'
}).reset_index()

# Merge aggregated transaction data back to customer profile
customer_profile = pd.merge(customers, customer_agg, on='CustomerID', how='left')

# Fill any missing values (e.g., customers with no transactions)
customer_profile.fillna({
    'TotalRevenue': 0,
    'TotalTransactions': 0,
    'AvgBasketSize': 0
}, inplace=True)

# Convert SignupDate to a datetime object for tenure calculation
customer_profile['SignupDate'] = pd.to_datetime(customer_profile['SignupDate'])

```

```
# Calculate tenure in days since signup
current_date = pd.Timestamp.now()
customer_profile['TenureDays'] = (current_date - customer_profile['SignupDate']).dt.days

# Drop unneeded columns for clustering
customer_clustering_data = customer_profile.drop(['CustomerName', 'SignupDate'], axis=1)
customer_clustering_data.head()
```

	CustomerID	Region	SignupMonth	TotalRevenue	TotalTransactions	AvgBasketSize	TenureDays
0	C0001	South America	2022-07	3354.52	5.0	2.400000	932
1	C0002	Asia	2022-02	1862.74	4.0	2.500000	1079
2	C0003	South America	2024-03	2725.38	4.0	3.500000	326
3	C0004	South America	2022-10	5354.88	8.0	2.875000	841
4	C0005	Asia	2022-08	2034.24	3.0	2.333333	896

Next steps: [Generate code with customer\\_clustering\\_data](#) [View recommended plots](#) [New interactive sheet](#)

## Feature Scaling:

```
from sklearn.preprocessing import StandardScaler

# Select numeric columns for clustering
numeric_columns = ['TotalRevenue', 'TotalTransactions', 'AvgBasketSize', 'TenureDays']
scaler = StandardScaler()
scaled_data = scaler.fit_transform(customer_clustering_data[numeric_columns])

# Convert scaled data back to a DataFrame for easier handling
scaled_df = pd.DataFrame(scaled_data, columns=numeric_columns)
scaled_df['CustomerID'] = customer_clustering_data['CustomerID']
scaled_df.head()
```

	TotalRevenue	TotalTransactions	AvgBasketSize	TenureDays	CustomerID
0	-0.051884	0.000000	-0.201382	1.152884	C0001
1	-0.862714	-0.451294	-0.030924	1.605593	C0002
2	-0.393842	-0.451294	1.673655	-0.713387	C0003
3	1.035375	1.353881	0.608293	0.872636	C0004
4	-0.769499	-0.902587	-0.315021	1.042017	C0005

Next steps: [Generate code with scaled\\_df](#) [View recommended plots](#) [New interactive sheet](#)

## 1. K-Means Clustering:

What is it?:

K-Means is a popular clustering algorithm that partitions data into K distinct clusters. Each customer is assigned to the cluster whose center (centroid) is nearest.

Why use it?:

Easy to understand and implement.

Works well when clusters are spherical or circular and data is relatively well-separated.

Good for large datasets.

```
#Step 1: Elbow Method to Find Optimal K

import matplotlib.pyplot as plt
from sklearn.cluster import KMeans

# Elbow method to find the optimal number of clusters
inertia = []


# Try different values of K from 1 to 10
for k in range(1, 11):
```

```

kmeans = KMeans(n_clusters=k, random_state=42)
kmeans.fit(scaled_df[numeric_columns]) # Fit the KMeans model on the scaled data
inertia.append(kmeans.inertia_) # Store inertia (sum of squared distances to the nearest cluster center)

# Plot the Elbow Method graph
plt.plot(range(1, 11), inertia, marker='o')
plt.title('Elbow Method')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('Inertia')
plt.show()

```



#Step 2: Apply K-Means with Chosen K:

```

# Apply K-Means clustering with K=4 (example)
kmeans = KMeans(n_clusters=4, random_state=42)
customer_profile['Cluster'] = kmeans.fit_predict(scaled_df[numeric_columns])

# Check the first few rows to see which customer belongs to which cluster
customer_profile.head()

```



	CustomerID	CustomerName	Region	SignupDate	SignupMonth	TotalRevenue	TotalTransactions	AvgBasketSize	TenureDays	Cluster
0	C0001	Lawrence Carroll	South America	2022-07-10	2022-07	3354.52	5.0	2.400000	932	
1	C0002	Elizabeth Lutz	Asia	2022-02-13	2022-02	1862.74	4.0	2.500000	1079	
2	C0003	Michael Rivera	South America	2024-03-07	2024-03	2725.38	4.0	3.500000	326	
		Kathleen	South							

Next steps:

[Generate code with customer\\_profile](#)
[View recommended plots](#)
[New interactive sheet](#)

#Step 3: Optional - Evaluate Clustering:

```

from sklearn.metrics import davies_bouldin_score

# Evaluate clustering quality
db_index = davies_bouldin_score(scaled_df[numeric_columns], customer_profile['Cluster'])
print('Davies-Bouldin Index:', db_index)

from sklearn.metrics import davies_bouldin_score

# Evaluate clustering quality
db_index = davies_bouldin_score(scaled_df[numeric_columns], customer_profile['Cluster'])
print('Davies-Bouldin Index:', db_index)

```



Davies-Bouldin Index: 1.215035501264994  
Davies-Bouldin Index: 1.215035501264994

#Step 4: Optional - Visualize the Clusters

```
from sklearn.decomposition import PCA
```

```
# Reduce the features to 2D using PCA
```

```
pca = PCA(n_components=2)
```

```
pca_components = pca.fit_transform(scaled_df[numeric_columns])
```

```
# Add PCA components to the DataFrame
```

```
customer_profile['PCA1'] = pca_components[:, 0]
```

```
customer_profile['PCA2'] = pca_components[:, 1]
```

```
# Plot the clusters
```

```
plt.scatter(customer_profile['PCA1'], customer_profile['PCA2'], c=customer_profile['Cluster'], cmap='viridis')
```


```
plt.title('Customer Segmentation (K-Means Clusters)')
```

```
plt.xlabel('PCA 1')
```

```
plt.ylabel('PCA 2')
```

```
plt.colorbar(label='Cluster')
```

```
plt.show()
```



#Step 5: Analyze the Clusters:

```
# Group by clusters and calculate the average for each feature
```

```
cluster_summary = customer_profile.groupby('Cluster').agg({
```

```
    'TotalRevenue': 'mean',
```

```
    'TotalTransactions': 'mean',
```

```
    'AvgBasketSize': 'mean',
```

```
    'TenureDays': 'mean'
```

```
}).reset_index()
```

```
# Print the cluster summary
```

```
print(cluster_summary)
```

Cluster	TotalRevenue	TotalTransactions	AvgBasketSize	TenureDays
0	5755.758125	7.708333	2.657543	709.958333
1	3476.467463	5.179104	2.505467	255.194030
2	2684.503091	3.709091	2.824545	775.872727
3	1104.939333	2.633333	1.761667	589.333333

#Step 6: Label the Clusters

```
# Create a new column for cluster labels based on the summary
```

```
def label_cluster(cluster_id):
```

```
    if cluster_id == 0:
```

```
        return 'High-Value, Frequent Shoppers'
```

```
    elif cluster_id == 1:
```

```
        return 'Moderate Shoppers'
```

```
    elif cluster_id == 2:
```

```
        return 'Loyal, Moderate Spenders'
```


```
    elif cluster_id == 3:
```

```
        return 'Low-Value, Infrequent Shoppers'
```

```
# Apply the label function to the Cluster column
```

```
customer_profile['ClusterLabel'] = customer_profile['Cluster'].apply(label_cluster)
```

```
# Check the updated customer profile
customer_profile.head()
```



	CustomerID	CustomerName	Region	SignupDate	SignupMonth	TotalRevenue	TotalTransactions	AvgBasketSize	TenureDays	Cluster
0	C0001	Lawrence Carroll	South America	2022-07-10	2022-07	3354.52	5.0	2.400000	932	
1	C0002	Elizabeth Lutz	Asia	2022-02-13	2022-02	1862.74	4.0	2.500000	1079	
2	C0003	Michael Rivera	South America	2024-03-07	2024-03	2725.38	4.0	3.500000	326	
3	C0004	Kathleen	South America	2022-10-09	2022-10	5354.88	8.0	2.875000	841	

Next steps:

[Generate code with customer\\_profile](#)

[View recommended plots](#)

[New interactive sheet](#)

Cluster:

- Cluster 0: High revenue, frequent transactions, and larger average basket size.
- Cluster 1: Moderate revenue, fewer transactions, shorter tenure.
- Cluster 2: Moderate revenue, larger basket size, longer tenure.
- Cluster 3: Low revenue, infrequent transactions, small basket size.

```
# Save the customer profile with cluster labels to a new CSV
customer_profile.to_csv('CustomerProfile_with_Clusters.csv', index=False)
```

Evaluate the Clustering Performance:

```
from sklearn.metrics import davies_bouldin_score
db_index = davies_bouldin_score(scaled_df[numeric_columns], kmeans.labels_)
print("DB Index:", db_index)
```



DB Index: 1.215035501264994

```
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
pca_components = pca.fit_transform(scaled_df[numeric_columns])

plt.scatter(pca_components[:, 0], pca_components[:, 1], c=kmeans.labels_, cmap='viridis')
plt.title("Customer Segmentation using K-means")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.show()
```

